



Problem Statement - Prepare a classification model using SVM for salary data ¶

1. Import Necessary Libraries

```
In [6]: import numpy as np
import pandas as pd
import tensorflow as tf
import seaborn as sns
```

2. Import Data

```
In [7]: Salary_data_train = pd.read_csv('SalaryData_Train(1).csv')
Salary_data_test = pd.read_csv('SalaryData_Test(1).csv')
Salary_data_train.columns
Salary_data_test.columns
String_columns = ['workclass', 'education', 'maritalstatus', 'occupation', 'relat
```

```
In [8]: Salary_data_test.columns
```

```
Out[8]: Index(['age', 'workclass', 'education', 'educationno', 'maritalstatus',
              'occupation', 'relationship', 'race', 'sex', 'capitalgain',
              'capitalloss', 'hoursperweek', 'native', 'Salary'],
              dtype='object')
```

In [9]: Salary_data_test

Out[9]:

	age	workclass	education	educationno	maritalstatus	occupation	relationship	race
0	25	Private	11th	7	Never-married	Machine-op-inspct	Own-child	Black
1	38	Private	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband	White
2	28	Local-gov	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husband	White
3	44	Private	Some-college	10	Married-civ-spouse	Machine-op-inspct	Husband	Black
4	34	Private	10th	6	Never-married	Other-service	Not-in-family	White
...
15055	33	Private	Bachelors	13	Never-married	Prof-specialty	Own-child	White
15056	39	Private	Bachelors	13	Divorced	Prof-specialty	Not-in-family	White
15057	38	Private	Bachelors	13	Married-civ-spouse	Prof-specialty	Husband	White
15058	44	Private	Bachelors	13	Divorced	Adm-clerical	Own-child	Asian-Pac-Islander
15059	35	Self-emp-inc	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White

15060 rows × 14 columns



```
In [10]: Salary_data_train
```

```
Out[10]:
```

	age	workclass	education	educationno	maritalstatus	occupation	relationship	race	sex
0	39	State-gov	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	M
1	50	Self-emp-not-inc	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	M
2	38	Private	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	M
3	53	Private	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	M
4	28	Private	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	F
...
30156	27	Private	Assoc-acdm	12	Married-civ-spouse	Tech-support	Wife	White	F
30157	40	Private	HS-grad	9	Married-civ-spouse	Machine-op-inspct	Husband	White	M
30158	58	Private	HS-grad	9	Widowed	Adm-clerical	Unmarried	White	F
30159	22	Private	HS-grad	9	Never-married	Adm-clerical	Own-child	White	M
30160	52	Self-emp-inc	HS-grad	9	Married-civ-spouse	Exec-managerial	Wife	White	F

30161 rows × 14 columns

```
In [11]: Salary_data_train.columns
```

```
Out[11]: Index(['age', 'workclass', 'education', 'educationno', 'maritalstatus',
               'occupation', 'relationship', 'race', 'sex', 'capitalgain',
               'capitalloss', 'hoursperweek', 'native', 'Salary'],
              dtype='object')
```

3. Data Understanding

```
In [12]: Salary_data_train.shape
```

```
Out[12]: (30161, 14)
```

```
In [13]: Salary_data_test.shape
```

```
Out[13]: (15060, 14)
```

```
In [14]: Salary_data_train.dtypes,Salary_data_test.dtypes
```

```
Out[14]: (age          int64
workclass    object
education    object
educationno   int64
maritalstatus object
occupation   object
relationship object
race         object
sex          object
capitalgain   int64
capitalloss   int64
hoursperweek  int64
native       object
Salary       object
dtype: object,
age          int64
workclass    object
education    object
educationno   int64
maritalstatus object
occupation   object
relationship object
race         object
sex          object
capitalgain   int64
capitalloss   int64
hoursperweek  int64
native       object
Salary       object
dtype: object)
```

```
In [15]: Salary_data_train.describe(),Salary_data_test.describe()
```

```
Out[15]: (
      count  30161.000000  30161.000000  30161.000000  30161.000000  30161.000000
      mean    38.438115    10.121316    1092.044064    88.302311    40.931269
      std     13.134830     2.550037    7406.466611    404.121321    11.980182
      min     17.000000     1.000000     0.000000     0.000000     1.000000
      25%     28.000000     9.000000     0.000000     0.000000    40.000000
      50%     37.000000    10.000000     0.000000     0.000000    40.000000
      75%     47.000000    13.000000     0.000000     0.000000    45.000000
      max     90.000000    16.000000   99999.000000   4356.000000   99.000000,
      count  15060.000000  15060.000000  15060.000000  15060.000000  15060.000000
      mean    38.768327    10.112749    1120.301594    89.041899    40.951594
      std     13.380676     2.558727    7703.181842    406.283245    12.062831
      min     17.000000     1.000000     0.000000     0.000000     1.000000
      25%     28.000000     9.000000     0.000000     0.000000    40.000000
      50%     37.000000    10.000000     0.000000     0.000000    40.000000
      75%     48.000000    13.000000     0.000000     0.000000    45.000000
      max     90.000000    16.000000   99999.000000   3770.000000   99.000000)
```



4. Data Preparation

```
In [16]: from sklearn.preprocessing import LabelEncoder
```

```
In [17]: LabelEncoder = LabelEncoder()
```

```
In [18]: for i in String_columns:
          Salary_data_train[i] = LabelEncoder.fit_transform(Salary_data_train[i])
          Salary_data_test[i] = LabelEncoder.fit_transform(Salary_data_test[i])
```

```
In [19]: Salary_data_train
```

Out[19]:

	age	workclass	education	educationno	maritalstatus	occupation	relationship	race	sex
0	39	5	9	13	4	0	1	4	1
1	50	4	9	13	2	3	0	4	1
2	38	2	11	9	0	5	1	4	1
3	53	2	1	7	2	5	0	2	1
4	28	2	9	13	2	9	5	2	0
...
30156	27	2	7	12	2	12	5	4	0
30157	40	2	11	9	2	6	0	4	1
30158	58	2	11	9	6	0	4	4	0
30159	22	2	11	9	4	0	3	4	1
30160	52	3	11	9	2	3	5	4	0

30161 rows × 14 columns



```
In [20]: Salary_data_test
```

Out[20]:

	age	workclass	education	educationno	maritalstatus	occupation	relationship	race	sex
0	25	2	1	7	4	6	3	2	1
1	38	2	11	9	2	4	0	4	1
2	28	1	7	12	2	10	0	4	1
3	44	2	15	10	2	6	0	2	1
4	34	2	0	6	4	7	1	4	1
...
15055	33	2	9	13	4	9	3	4	1
15056	39	2	9	13	0	9	1	4	0
15057	38	2	9	13	2	9	0	4	1
15058	44	2	9	13	0	0	3	1	1
15059	35	3	9	13	2	3	0	4	1

15060 rows × 14 columns



```
In [21]: Salary_data_train.dtypes,Salary_data_test.dtypes
```

```
Out[21]: (age                int64
workclass            int32
education            int32
educationno          int64
maritalstatus        int32
occupation            int32
relationship          int32
race                 int32
sex                  int32
capitalgain           int64
capitalloss           int64
hoursperweek          int64
native                int32
Salary               int32
dtype: object,
age                int64
workclass            int32
education            int32
educationno          int64
maritalstatus        int32
occupation            int32
relationship          int32
race                 int32
sex                  int32
capitalgain           int64
capitalloss           int64
hoursperweek          int64
native                int32
Salary               int32
dtype: object)
```

5. Model Building

```
In [22]: X_train=Salary_data_train.iloc[0:500,0:13]
y_train=Salary_data_train.iloc[0:500,13]
X_test=Salary_data_test.iloc[0:300,0:13]
y_test=Salary_data_test.iloc[0:300,13]
```

```
In [23]: X_train
```

Out[23]:

	age	workclass	education	educationno	maritalstatus	occupation	relationship	race	sex	ci
0	39	5	9	13	4	0	1	4	1	
1	50	4	9	13	2	3	0	4	1	
2	38	2	11	9	0	5	1	4	1	
3	53	2	1	7	2	5	0	2	1	
4	28	2	9	13	2	9	5	2	0	
...
495	31	3	9	13	4	4	1	4	0	
496	44	2	15	10	2	13	0	4	1	
497	29	2	10	16	4	9	3	4	1	
498	30	2	15	10	4	5	4	4	0	
499	27	2	11	9	4	13	1	4	1	

500 rows × 13 columns



```
In [24]: y_train
```

Out[24]:

0	0
1	0
2	0
3	0
4	0
...	..
495	0
496	0
497	0
498	0
499	0

Name: Salary, Length: 500, dtype: int32

In [25]: X_test

Out[25]:

	age	workclass	education	educationno	maritalstatus	occupation	relationship	race	sex	ci
0	25	2	1	7	4	6	3	2	1	
1	38	2	11	9	2	4	0	4	1	
2	28	1	7	12	2	10	0	4	1	
3	44	2	15	10	2	6	0	2	1	
4	34	2	0	6	4	7	1	4	1	
...
295	56	3	9	13	2	3	0	4	1	
296	37	2	11	9	2	2	0	4	1	
297	52	2	3	2	2	6	0	4	1	
298	26	2	11	9	5	5	3	4	1	
299	33	2	12	14	2	9	0	4	1	

300 rows × 13 columns

In [26]: y_test

Out[26]:

0	0
1	0
2	1
3	1
4	0
...	..
295	0
296	0
297	0
298	0
299	0

Name: Salary, Length: 300, dtype: int32

6. Model Training, Testing & Evaluation

In [35]: `from sklearn.svm import SVC`

```
In [36]: model_linear =SVC(kernel = 'linear')
model_linear.fit(X_train,y_train)
train_pred_lin = model_linear.predict(X_train)
test_pred_lin = model_linear.predict(X_test)
train_lin_acc = np.mean(train_pred_lin==y_train)
test_lin_acc= np.mean(test_pred_lin==y_test)
print("Accuracy Of train-data using Linear", train_lin_acc)
print("Accuracy of test-data using Linear",test_lin_acc )
```

Accuracy Of train-data using Linear 0.818
Accuracy of test-data using Linear 0.8166666666666667

```
In [37]: model_poly=SVC(kernel='poly')
model_poly.fit(X_train,y_train)
train_pred_poly=model_poly.predict(X_train)
test_pred_poly=model_poly.predict(X_test)
train_poly_acc=np.mean(train_pred_poly==y_train)
test_poly_acc=np.mean(test_pred_poly==y_test)
print("Accuracy of train-data using POLY",train_poly_acc)
print("Accuracy of test-data using POLY",test_poly_acc)
```

Accuracy of train-data using POLY 0.812
Accuracy of test-data using POLY 0.8033333333333333

```
In [38]: model_rbf=SVC(kernel='rbf')
model_rbf.fit(X_train,y_train)
train_pred_rbf=model_rbf.predict(X_train)
test_pred_rbf=model_rbf.predict(X_test)
train_rbf_acc=np.mean(train_pred_rbf==y_train)
test_rbf_acc=np.mean(test_pred_rbf==y_test)
print("Accuracy of train-data using rbf",train_rbf_acc)
print("Accuracy of test-data using rbf",test_rbf_acc)
```

Accuracy of train-data using rbf 0.812
Accuracy of test-data using rbf 0.8033333333333333

Conclusion: The accuracy of model is good with linear kernel tric